

Package ‘forestError’

July 29, 2019

Type Package

Title Individualized Prediction Errors and Intervals for Random Forests

Version 0.0.0.9000

Author Benjamin Lu and Johanna Hardin

Maintainer Benjamin Lu <b.lu@berkeley.edu>

Description Estimates both conditional mean squared prediction errors and conditional prediction intervals for random forest predictions. The prediction errors and intervals are conditional in the sense that each error and interval is specific to the individual observation whose response is being predicted. More details can be found in Lu and Hardin (2019+) (in preparation).

Imports Rcpp

License GPL-3

Encoding UTF-8

LazyData true

RoxygenNote 6.1.1

LinkingTo Rcpp

R topics documented:

perror	1
qerror	2
quantForestError	3
Index	6

perror	<i>Empirical cumulative distribution functions of conditional errors</i>
--------	--

Description

Returns probabilities from the estimated cumulative distribution function of the conditional error distribution associated with each test prediction.

Usage

```
perror(p, xs = 1:n.test)
```

Arguments

p	Vector of quantiles
xs	Indices of the test observations for which the conditional error CDFs are desired. Defaults to all test observations given in the call of quantForestError.

Value

If either p or xs has length one, then a vector is returned with the desired probabilities. If both have length greater than one, then a data.frame is returned, with rows corresponding to the inputted xs and columns corresponding to the inputted p.

Author(s)

Benjamin Lu <b.lu@berkeley.edu>; Johanna Hardin <jo.hardin@pomona.edu>

Examples

```
# get the probability that the error associated with each test prediction
# is less than -5 and the probability that the error associated with each
# test prediction is less than 5
peror(c(-5, 5))

# same as above but only for the first three test observations
peror(c(-5, 5), 1:3)
```

qerror

Empirical quantile functions of conditional errors

Description

Returns quantiles from the estimated quantile function of the conditional error distribution associated with each test prediction.

Usage

```
qerror(q, xs = 1:n.test)
```

Arguments

q	Vector of probabilities
xs	Indices of the test observations for which the conditional error CDFs are desired. Defaults to all test observations given in the call of quantForestError.

Value

If either `q` or `xs` has length one, then a vector is returned with the desired probabilities. If both have length greater than one, then a `data.frame` is returned, with rows corresponding to the inputted `xs` and columns corresponding to the inputted `q`.

Author(s)

Benjamin Lu <b.lu@berkeley.edu>; Johanna Hardin <jo.hardin@pomona.edu>

Examples

```
# get the 0.25 and 0.8 quantiles of the error distribution associated with
# each test observation
qerror(c(0.25, 0.8))

# same as above but only for the first three test observations
qerror(c(0.25, 0.8), 1:3)
```

quantForestError	<i>Quantify random forest prediction error</i>
------------------	--

Description

Estimates the conditional mean squared prediction errors, conditional prediction intervals, and conditional error distributions of random forest predictions.

Usage

```
quantForestError(forest, X.train, X.test, Y.train = NULL,
  what = c("mspe", "interval", "cons.interval", "p.error", "q.error"),
  alpha = 0.05, rcpp = TRUE)
```

Arguments

<code>forest</code>	The random forest object being used for prediction.
<code>X.train</code>	A matrix or <code>data.frame</code> with the observations that were used to train forest; each row should be an observation, and each column should be a predictor variable.
<code>X.test</code>	A matrix or <code>data.frame</code> with the observations to be predicted; each row should be an observation, and each column should be a predictor variable.
<code>Y.train</code>	A vector of the responses of the observations that were used to train forest. Required if forest was created using the <code>ranger</code> package, but not if forest was created using the <code>randomForest</code> package or the <code>randomForestSRC</code> package.
<code>what</code>	A vector of characters indicating what outputs are desired. Possible options are conditional mean squared prediction error estimates (<code>"mspe"</code>), conditional prediction intervals (<code>"interval"</code>), conservative conditional prediction intervals (<code>"cons.interval"</code>), the conditional error distribution functions (<code>"p.error"</code>), and the conditional quantile functions (<code>"q.error"</code>).

alpha	The type-I error rate desired for the conditional prediction intervals; required if "interval" or "cons.interval" are included in what.
rcpp	A logical indicating whether the weights should be computed in C++ using the Rcpp package for reduced runtime. Recommended especially when the number of training observations, test observations, or trees is large. Defaults to TRUE.

Details

When training the random forest using `randomForest`, `ranger` or `quantregForest`, `keep.inbag` must be set to TRUE. When training the random forest using `randomForestSRC`, `membership` must be set to TRUE.

The random forest predictions are always returned as a dataframe. Additional columns are included in the dataframe depending on the user's selections in the argument `what`. In particular, including `"mspe"` in `what` will add an additional column with the conditional mean squared prediction error of each test prediction to the dataframe; including `"interval"` in `what` will add two additional columns with the lower and upper bounds of a conditional prediction interval for each test prediction to the dataframe; and including `"cons.interval"` will add to the dataframe two additional columns with the lower and upper bounds of a conservative prediction interval for each test prediction. Conservative prediction intervals are generated by breaking ties in the empirical error distribution conservatively. They may be desirable when the number of observations is small.

If `"p.error"` or `"q.error"` is included in `what`, then a list will be returned as output. The first element of the list, named `"estimates"`, is the dataframe described in the above paragraph. The other one or two elements of the list are the empirical CDFs (`perror`) and/or the empirical quantile functions (`qerror`) of the conditional error distributions associated with the test predictions.

Value

A dataframe with one or more of the following columns, as described in the details section:

pred	The random forest predictions of the test observations
error	The estimated conditional mean square prediction errors of the random forest predictions
lower	The estimated lower bounds of the conditional prediction intervals for the test observations
upper	The estimated upper bounds of the conditional prediction intervals for the test observations
lowerCons	The conservatively estimated lower bounds of the conditional prediction intervals for the test observations
upperCons	The conservatively estimated upper bounds of the conditional prediction intervals for the test observations

In addition, one or both of the following functions, as described in the details section:

perror	The empirical cumulative distribution functions of the conditional error distributions associated with the test predictions
qerror	The empirical quantile functions of the conditional error distributions associated with the test predictions

Author(s)

Benjamin Lu <b.lu@berkeley.edu>; Johanna Hardin <jo.hardin@pomona.edu>

Examples

[illegible]

Index

perror, [1](#)

qerror, [2](#)

quantForestError, [3](#)